

Support Vector Machines: Methods and Applications

Exercise Session II

Ricardo Castro Lynn Houthuys Zahra Karevan

{ricardo.castro,lynn.houthuys,zahra.karevan}@esat.kuleuven.be

Introduction

The Matlab scripts, related documents, referred papers and data-sets are available for academic purposes on Toledo: <http://toledo.kuleuven.be/>. The toolbox to be used for LS-SVM should be downloaded from <http://www.esat.kuleuven.be/sista/lssvmlab/>.

2 Exercise Session 2: Function Estimation and Time-series Prediction

2.1 The Support Vector Machine for Regression

For this exercise the SVM toolbox is used. Please download it directly from Toledo: [Toledo website](#) → [SVM Exercises Course](#) → [Course Documents](#) → [SVM toolbox](#).

Execute:

```
>> uiregress
```

Construct a data-set of around 20 data-points. Which kernel is best suited for your data-set? Construct a data-set where a linear kernel is better than any other kernel. What is the influence of different values of ϵ ? (try small values *e.g.*, 0.1, 0.25, 0.5...) and of Bound (*i.e.* hyperparameter C in the previous exercise session)? Try values with larger increments *e.g.* 0.01, 0.1, 1, 10, 100. Where does the sparsity property come in? In what respect is SVM regression different from a classical Least Squares fit?

2.2 A Simple Example: The *sinc*

For the remainder of this exercise session, LS-SVMlab is used ¹.

Run the demo demofun for a demonstration. We start with the following synthetic example:

```
>> X = (-3:0.01:3)';  
>> Y = sinc(X)+0.1.*randn(length(X), 1);
```

The training/validation and test sets are created:

```
>> Xtrain = X(1:2:length(X));  
>> Ytrain = Y(1:2:length(Y));  
>> Xtest = X(2:2:length(X));  
>> Ytest = Y(2:2:length(Y));
```

The following commands can be used to make an LS-SVM model with the RBF kernel and arbitrary values of the hyper-parameters:

¹Be sure to be using the latest version of the LS-SVMlab toolbox. You can download it from <http://www.esat.kuleuven.be/sista/lssvmlab/>

```
>> gam = 100;
>> sig2 = 1.0;
>> [alpha,b] = trainlssvm({Xtrain,Ytrain,'f',gam,sig2,'RBF_kernel'});
```

The results on the training set can be visualized via:

```
>> plotlssvm({Xtrain,Ytrain,'f',gam,sig2,'RBF_kernel'},{alpha,b});
```

The results of applying the trained model on the test set are calculated using:

```
>> YtestEst = simlssvm({Xtrain,Ytrain,'f',gam,sig2,'RBF_kernel'},...
    {alpha,b},Xtest);
```

and visualized with:

```
>> plot(Xtest,Ytest, '.');
>> hold on;
>> plot(Xtest,YtestEst, 'r+');
>> legend('Ytest','YtestEst');
```

Make similar plots showing the results on training and test sets with different values of `gam` and `sig2`. Do you think there is one pair of optimal hyper-parameters?

2.3 Hyper-parameter Tuning

The functions `crossvalidate` and `leaveoneout` can be used to estimate the performance of a set of hyper-parameters:

```
>> cost_crossval = crossvalidate({Xtrain,Ytrain,'f',gam,sig2},10);
>> cost_loo = leaveoneout({Xtrain,Ytrain,'f',gam,sig2});
```

The optimization of the hyper-parameters is done with the function `tunelssvm`. For example,

```
>> optFun = 'gridsearch';
>> [gam,sig2,cost] = tunelssvm({X,Y,'f',[],[],'RBF_kernel'}, ...
    optFun,'crossvalidate_lssvm',{10,'mse'});
```

tunes the hyper-parameters using 10-fold crossvalidation using a gridsearch approach and Coupled Simulated Annealing. Train a model with tuned hyper-parameters and display the results on the training data:

```
>> [alpha,b] = trainlssvm({Xtrain,Ytrain,'f',gam,sig2});
>> plotlssvm({Xtrain,Ytrain,'f',gam,sig2},{alpha,b});
```

Run the previous `tunelssvm` command several times. What can you say about the values of the hyperparameters and the results? Change `optFun` to `'simplex'` (`>> optFun = 'simplex';`) and run the tuning once more. What is the difference between `'simplex'` and `'gridsearch'`?

2.4 Application of the Bayesian Framework

Let's consider the same toy example. The Bayesian framework now can be used to tune and to analyze the LS-SVM regressor. The basic result from the Bayesian framework for the LS-SVM is the derivation of *the probability that the data points are generated by the given model*. This is called the posterior probability. This probability criterion is expressed as a number. There are 3 variants: the posterior with respect to the model parameters `alpha` and `b`, the posterior with respect to the regularization constant `gam` and the posterior with respect to the choice of the kernel and its parameter. The cost (negative logarithm of the posteriors) is calculated by the function call

```
>> sig2 = 0.5; gam = 10;
>> criterion_L1 = bay_lssvm({Xtrain,Ytrain,'f',gam,sig2},1);
>> criterion_L2 = bay_lssvm({Xtrain,Ytrain,'f',gam,sig2},2);
>> criterion_L3 = bay_lssvm({Xtrain,Ytrain,'f',gam,sig2},3);
```

The model can be optimized with respect to these criteria:

```
>> [~,alpha,b] = bay_optimize({Xtrain,Ytrain,'f',gam,sig2},1);
>> [~,gam] = bay_optimize({Xtrain,Ytrain,'f',gam,sig2},2);
>> [~,sig2] = bay_optimize({Xtrain,Ytrain,'f',gam,sig2},3);
```

Think of a clear schematic visualization of this three-level principle.

For regression, the error bars can be computed using Bayesian inference, *e.g.*

```
>> sig2e = bay_errorbar({Xtrain,Ytrain,'f',gam,sig2},'figure');
```

For classification, it is also possible to get probability estimates. Load the Iris data-set:

```
>> clear;
>> load iris;
>> gam = 5; sig2 = 0.75;
```

The probabilities that a certain data point belongs to the positive class given the model is calculated by:

```
>> bay_modoutClass({X,Y,'c',gam,sig2},'figure');
```

How do you interpret the colors? (hint: activate the color-bar of the figure). Change the values of `gam` and `sig2`. What is the influence of the given hyper-parameters on this figure?

The Bayesian framework can also be used to select the most relevant inputs by Automatic Relevance Determination (ARD). The following procedure uses this criterion for *backward selection* for a three dimensional input selection task:

```
>> X = 6.*rand(100,3)-3;
>> Y = sinc(X(:,1)) + 0.1.*randn(100,1);
>> [selected, ranking] = bay_lssvmARD({X,Y,'class',gam,sig2});
```

Can you visualize the results on a simple figure? Think of a similar way to do input selection by using the `crossvalidate` function instead of the Bayesian framework?

2.5 Robust Regression

In situations where the data is corrupted with non-Gaussian noise or outliers, it becomes important to incorporate robustness into the estimation. Consider the following simple example:

```
>> X = (-6:0.2:6)';
>> Y = sinc(X) + 0.1.*rand(size(X));
```

Outliers are added via:

```
>> out = [15 17 19];
>> Y(out) = 0.7+0.3*rand(size(out));
>> out = [41 44 46];
>> Y(out) = 1.5+0.2*rand(size(out));
```

Train and plot an LS-SVM regression model with `gam = 100`, `sig2 = 0.1`. What is the influence of the outlying data points?

Let's now train a robust LS-SVM model using the object-oriented interface and robust crossvalidation.




Figure 1: Schematic representation of the nonlinear auto-regressive model.

```
>> model = initlssvm(X,Y,'f',[[],[]],'RBF_kernel');
>> costFun = 'rcrossvalidate_lssvm';
>> wFun = 'whuber';
>> model = tunelssvm(model,'simplex',costFun,{10,'mae'},wFun);
>> model = robustlssvm(model);
>> plotlssvm(model);
```

Discuss the results. Why in this case is the *mean absolute error* (mae) preferred over the classical *mean squared error* (mse)? Try alternatives to the weighting function wFun (e.g. 'whampel', 'wlogistic', 'wmyriad'). Check the user's guide for more information.

2.6 Homework Problem

2.6.1 Introduction: Time-series Prediction

The linear auto-regressive (AR) model of a process Z_t with $t = 1, 2, \dots, \infty$ is given by

$$\hat{z}_t = a_1 z_{t-1} + a_2 z_{t-2} + \dots + a_n z_{t-n} , \quad (1)$$

with $a_i \in \mathbb{R}$ for $i = 1, \dots, n$ and n the model order. At the same time, the nonlinear variant (NAR) is described as:

$$\hat{z}_t = f(z_{t-1}, z_{t-2}, \dots, z_{t-n}) . \quad (2)$$

A depiction of these processes can be found in Figure 1. Remark that in this way, *the time-series identification can be written as a classical black-box regression modeling problem*:

$$\hat{y}_t = f(x_t) , \quad (3)$$

with $y_t = z_t$ and $x_t = [z_{t-1}, z_{t-2}, \dots, z_{t-n}]$. In a nutshell, any provided sequence Z can be mapped into a regression problem. In our case, the required transformation can be done via e.g. the command `windowize`:

```
>> load logmap;
>> order = 10;
>> X = windowize(Z,1:(order+1));
>> Y = X(:,end);
>> X = X(:,1:order);
```

Now a model can be built using these data points:

```
>> gam = 10; sig2 = 10;
>> [alpha,b] = trainlssvm({X,Y,'f',gam,sig2});
```

As in the previous section, the `gam` and `sig2` can be optimized using `crossvalidate`. In the same way, one can optimize `order` as a parameter. It is straightforward to predict the next time-step with this model:

```
>> Xnew = Z((end-order+1):end)';
>> Z(end+1) = simlssvm({X,Y,'f',gam,sig2},{alpha,b},Xnew);
```

Note that this procedure can be *recurrently applied* by employing the following call:

```
>> horizon = length(Ztest)-order;
>> Zpt = predict({X,Y,'f',gam,sig2},Ztest(1:order),horizon);
>> plot([Ztest(order+1:end) Zpt]);
```

Here we obtained our predictions *w.r.t.* the test set `Ztest` which can be constructed by splitting the initial sequence (`Z`) into two subsequent counterparts, for instance:

```
>> test_size = 100;
>> Ztrain = Z(1:length(Z)-test_size);
>> Ztest = Z(length(Z)-test_size+1:end);
```

What are the observed results? Try to change the order of the model and evaluate the mse (mean square error) on the test set. How does the order of the model affect the obtained performance?

2.6.2 Application: Santa Fe Laser Dataset

We would like you to apply now this procedure to the the Santa Fe laser time series data-set and answer the following set of questions. **Please, justify all of your answers as thoroughly as possible.** Keep in mind that one of the skills being evaluated during the final examination is also your ability to creatively and constructively address problems with which you may not be entirely familiar with the help of the tools learned during the course.

Before proceeding with the experiments on the Santa Fe Laser dataset ensure that your data is already split to the training (`Z`) and test (`Ztest`) sets as it was described in the previous section.

Questions:

- Does `order=50` for the utilized auto-regressive model sound like a good choice?
- Would it be sensible to use the performance of this recurrent prediction on the *validation set* to optimize the *hyper-parameters* and the *model order*?